

TASK 9

September 15, 2025

```
[1]: import pandas as pd
      from glob import glob
```

```
[4]: import os
      print(os.getcwd())
```

/home/01e8c0d6-55b3-4404-abea-629433cfc2c8

```
[7]: %pip install pandas numpy matplotlib seaborn scikit-learn plotly missingno
      ↪--quiet

import os, glob
print("Working directory:", os.getcwd())
print("CSV files found here:")
for f in sorted(glob.glob("*.csv")):
    print("  ", f)
```

Note: you may need to restart the kernel to use updated packages.

Working directory: /home/01e8c0d6-55b3-4404-abea-629433cfc2c8

CSV files found here:

```
gender.csv
olist_customers_dataset.csv
olist_geolocation_dataset.csv
olist_order_items_dataset.csv
olist_order_payments_dataset.csv
olist_order_reviews_dataset.csv
olist_orders_dataset.csv
olist_products_dataset.csv
olist_sellers_dataset.csv
onlineretail.csv
product_category_name_translation.csv
test.csv
train.csv
```

```
[8]: import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns
```

```
import glob, os
```

```
sns.set_style("whitegrid")
```

```
pd.set_option('display.max_columns', 80)
```

```
pd.set_option('display.width', 120)
```

```
[9]: orders = pd.read_csv("olist_orders_dataset.csv", low_memory=False)
order_items = pd.read_csv("olist_order_items_dataset.csv", low_memory=False)
products = pd.read_csv("olist_products_dataset.csv", low_memory=False)
customers = pd.read_csv("olist_customers_dataset.csv", low_memory=False)
sellers = pd.read_csv("olist_sellers_dataset.csv", low_memory=False)
payments = pd.read_csv("olist_order_payments_dataset.csv", low_memory=False)
reviews = pd.read_csv("olist_order_reviews_dataset.csv", low_memory=False)

if os.path.exists("olist_geolocation_dataset.csv"):
    geolocation = pd.read_csv("olist_geolocation_dataset.csv", low_memory=False)
if os.path.exists("product_category_name_translation.csv"):
    category_translation = pd.read_csv("product_category_name_translation.csv",
    ↪low_memory=False)

print("orders:", orders.shape, "order_items:", order_items.shape, "products:",
    ↪products.shape)
orders.head(2)
```

```
orders: (99441, 8) order_items: (112650, 7) products: (32951, 9)
```

```
[9]:
```

	order_id	customer_id	
order_status	order_purchase_timestamp \		
0	e481f51cbdc54678b7cc49136f2d6af7	9ef432eb6251297304e76186b10a928d	
delivered	2017-10-02 10:56:33		
1	53cdb2fbc8bc7dce0b6741e2150273451	b0830fb4747a6c6d20dea0b8c802d7ef	
delivered	2018-07-24 20:41:37		
	order_approved_at	order_delivered_carrier_date	
order_delivered_customer_date	order_estimated_delivery_date		
0	2017-10-02 11:07:15	2017-10-04 19:55:00	2017-10-10
21:25:13	2017-10-18 00:00:00		
1	2018-07-26 03:24:27	2018-07-26 14:31:00	2018-08-07
15:27:45	2018-08-13 00:00:00		

```
[10]: print(orders.info())
print(order_items.info())
print(products.columns.tolist())

date_cols =
    ↪["order_purchase_timestamp", "order_approved_at", "order_delivered_carrier_date",
```

```

        "order_delivered_customer_date", "order_estimated_delivery_date"]
for c in date_cols:
    if c in orders.columns:
        orders[c] = pd.to_datetime(orders[c], errors="coerce")

if 'order_status' in orders.columns:
    print(orders['order_status'].value_counts())

```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 99441 entries, 0 to 99440
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
0	order_id	99441 non-null	object
1	customer_id	99441 non-null	object
2	order_status	99441 non-null	object
3	order_purchase_timestamp	99441 non-null	object
4	order_approved_at	99281 non-null	object
5	order_delivered_carrier_date	97658 non-null	object
6	order_delivered_customer_date	96476 non-null	object
7	order_estimated_delivery_date	99441 non-null	object

```
dtypes: object(8)
```

```
memory usage: 6.1+ MB
```

```
None
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 112650 entries, 0 to 112649
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	order_id	112650 non-null	object
1	order_item_id	112650 non-null	int64
2	product_id	112650 non-null	object
3	seller_id	112650 non-null	object
4	shipping_limit_date	112650 non-null	object
5	price	112650 non-null	float64
6	freight_value	112650 non-null	float64

```
dtypes: float64(2), int64(1), object(4)
```

```
memory usage: 6.0+ MB
```

```
None
```

```

['product_id', 'product_category_name', 'product_name_lenght',
'product_description_lenght', 'product_photos_qty', 'product_weight_g',
'product_length_cm', 'product_height_cm', 'product_width_cm']

```

```
order_status
```

```
delivered      96478
```

```
shipped        1107
```

```
canceled        625
```

```
unavailable     609
```

```

invoiced      314
processing    301
created        5
approved      2
Name: count, dtype: int64

```

```

[11]: orders.columns = orders.columns.str.strip().str.lower()
      order_items.columns = order_items.columns.str.strip().str.lower()
      products.columns = products.columns.str.strip().str.lower()
      customers.columns = customers.columns.str.strip().str.lower()
      payments.columns = payments.columns.str.strip().str.lower()
      reviews.columns = reviews.columns.str.strip().str.lower()
      sellers.columns = sellers.columns.str.strip().str.lower()

      orders = orders.drop_duplicates()
      order_items = order_items.drop_duplicates()

      print("Missing in orders:\n", orders.isnull().sum().
            ↪sort_values(ascending=False).head(8))

```

```

Missing in orders:
  order_delivered_customer_date    2965
  order_delivered_carrier_date     1783
  order_approved_at                160
  order_id                         0
  order_purchase_timestamp          0
  order_status                     0
  customer_id                      0
  order_estimated_delivery_date     0
dtype: int64

```

```

[12]: order_totals = order_items.groupby("order_id").agg(
      n_items = ("order_item_id", "count"),
      items_value = ("price", "sum"),
      freight_value = ("freight_value", "sum")
    ).reset_index()
      order_totals["order_revenue"] = order_totals["items_value"] +
      ↪order_totals["freight_value"]

      payments_by_order = payments.groupby("order_id").agg(
        payment_count = ("payment_type", "count"),
        payment_value_sum = ("payment_value", "sum"),
        payment_types = ("payment_type", lambda x: ", ".join(x.unique()))
      ).reset_index()

```

```

reviews_by_order = reviews.groupby("order_id").agg(
    review_score = ("review_score", "mean"),
    review_count = ("review_comment_message", "count")
).reset_index()

df = orders.merge(order_totals, on="order_id", how="left")
df = df.merge(customers, on="customer_id", how="left")
df = df.merge(payments_by_order, on="order_id", how="left")
df = df.merge(reviews_by_order, on="order_id", how="left")

df[['order_id', 'order_purchase_timestamp', 'order_status', 'order_revenue', 'payment_value_sum',
    ↪head()

```

```

[12]:
order_id order_purchase_timestamp order_status
order_revenue payment_value_sum \
0 e481f51cbdc54678b7cc49136f2d6af7 2017-10-02 10:56:33 delivered
38.71 38.71
1 53cdb2fc8bc7dce0b6741e2150273451 2018-07-24 20:41:37 delivered
141.46 141.46
2 47770eb9100c2d0c44946d9cf07ec65d 2018-08-08 08:38:49 delivered
179.12 179.12
3 949d5b44dbf5de918fe9c16f97b45f8a 2017-11-18 19:28:06 delivered
72.20 72.20
4 ad21c59c0840e6cb83a9ceb5573f8159 2018-02-13 21:18:39 delivered
28.62 28.62

review_score
0 4.0
1 4.0
2 5.0
3 5.0
4 5.0

```

```

[13]: if 'order_delivered_customer_date' in df.columns and 'order_purchase_timestamp'
    ↪in df.columns:
        df['delivery_time_days'] = (df['order_delivered_customer_date'] -
    ↪df['order_purchase_timestamp']).dt.days

if 'order_approved_at' in df.columns:
    df['days_to_approve'] = (df['order_approved_at'] -
    ↪df['order_purchase_timestamp']).dt.days

```

```

if 'order_status' in df.columns and 'delivered' in df['order_status'].unique():
    df_sales = df[df['order_status'] == 'delivered'].copy()
else:
    df_sales = df.copy()

df_sales['order_month'] = df_sales['order_purchase_timestamp'].dt.to_period('M')
df_sales.head(2)

```

```

[13]:
order_id customer_id
order_status order_purchase_timestamp \
0 e481f51cbdc54678b7cc49136f2d6af7 9ef432eb6251297304e76186b10a928d
delivered 2017-10-02 10:56:33
1 53cdb2fc8bc7dce0b6741e2150273451 b0830fb4747a6c6d20dea0b8c802d7ef
delivered 2018-07-24 20:41:37

order_approved_at order_delivered_carrier_date order_delivered_customer_date
order_estimated_delivery_date \
0 2017-10-02 11:07:15 2017-10-04 19:55:00 2017-10-10 21:25:13
2017-10-18
1 2018-07-26 03:24:27 2018-07-26 14:31:00 2018-08-07 15:27:45
2018-08-13

n_items items_value freight_value order_revenue
customer_unique_id customer_zip_code_prefix \
0 1.0 29.99 8.72 38.71
7c396fd4830fd04220f754e42b4e5bff 3149
1 1.0 118.70 22.76 141.46
af07308b275d755c9edb36a90c618231 47813

customer_city customer_state payment_count payment_value_sum
payment_types review_score review_count \
0 sao paulo SP 3.0 38.71 credit_card,
voucher 4.0 1.0
1 barreiras BA 1.0 141.46
boleto 4.0 1.0

delivery_time_days days_to_approve order_month
0 8.0 0.0 2017-10
1 13.0 1.0 2018-07

```

```


[14]: oi = order_items.merge(products[['product_id', 'product_category_name']],
    on='product_id', how='left')
top_products = oi.groupby("product_id").agg(revenue=("price", "sum"),
    n_orders=("order_id", "nunique")).reset_index()
top_products = top_products.sort_values("revenue", ascending=False).head(10)
top_products

```

```
[14]:
```

	product_id	revenue	n_orders
24086	bb50f2e236e5eea0100680137654686c	63885.00	187
14068	6cdd53843498f92890544667809f1595	54730.20	151
27613	d6160fb7873f184099d9bc95e30376af	48899.34	35
27039	d1c427060a0f73f6b889a5c7c61f2ac4	47214.51	323
19742	99a4788cb24856965c36a24e339b6058	43025.56	467
8051	3dd2a17168ec895c781a9191c1e95ad7	41082.60	255
4996	25c38557cf793876c5abdd5931f922db	38907.32	38
12351	5f504b3a1c75b73d6151be81eb05bdc9	37733.90	63
10867	53b36df67ebb7c41585e8d54d6772e08	37683.42	306
22112	aca2eb7d00ea1a7b8ebd4e68314663af	37608.90	431

```
[15]: monthly = df_sales.groupby(df_sales['order_purchase_timestamp'].dt.
    ↳to_period('M')).agg(month_revenue=("order_revenue", "sum")).reset_index()
monthly['order_purchase_timestamp'] = monthly['order_purchase_timestamp'].dt.
    ↳to_timestamp()
plt.figure(figsize=(10,4))
plt.plot(monthly['order_purchase_timestamp'], monthly['month_revenue'])
plt.title("Monthly Revenue")
plt.ylabel("Revenue")
plt.xlabel("Month")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[16]: cust_orders = df_sales.groupby("customer_id").agg(
    n_orders=("order_id", "nunique"),
    lifetime_value=("order_revenue", "sum")
).reset_index()
repeat_rate = (cust_orders['n_orders'] > 1).mean()
print(f"Repeat purchase rate: {repeat_rate:.2%}")
print("Average Order Value (AOV):", df_sales['order_revenue'].mean())
```

Repeat purchase rate: 0.00%
Average Order Value (AOV): 159.82683876116837

```
[17]: seller_revenue = order_items.  
      ↪merge(products[['product_id', 'product_category_name']], on='product_id',  
      ↪how='left')  
seller_revenue = seller_revenue.groupby("seller_id").  
      ↪agg(revenue=("price", "sum")).reset_index().sort_values("revenue",  
      ↪ascending=False).head(10)  
seller_revenue
```

```
[17]:
```

	seller_id	revenue
857	4869f7a5dfa277a7dca6462dcf3b52b2	229472.63
1013	53243585a1d6dc2643021fd1853d8905	222776.05
881	4a3ca9315b744ce9f8e9374361493884	200472.92
3024	fa1c13f2614d7b5c4749cbc52fecda94	194042.03
1535	7c67e1448b00f6e969d365cea6b010ab	187923.89
1560	7e93a43ef30c4f03f38b393420bc753a	176431.87
2643	da8622b14eb17ae2831f4ac5b9dab84a	160236.57
1505	7a67c85e85bb2ce8582c35f2203ad736	141745.53
192	1025f0e2d44d7041d6cf58b6550e0bfa	138968.55
1824	955fee9216a65b617aa5c0531780ce60	135171.70

```
[20]: snapshot_date = df_sales['order_purchase_timestamp'].max() + pd.  
      ↪Timedelta(days=1)  
  
rfm = df_sales.groupby("customer_id").agg(  
      recency_days = ("order_purchase_timestamp", lambda x: (snapshot_date - x.  
      ↪max()).days),  
      frequency = ("order_id", "nunique"),  
      monetary = ("order_revenue", "sum")  
) .reset_index()  
  
# Create simple quartile scores (1-4)  
rfm['r_quartile'] = pd.qcut(rfm['recency_days'], 4, labels=[4,3,2,1]).  
      ↪astype(int) # smaller recency -> higher score  
rfm['f_quartile'] = pd.qcut(rfm['frequency'].rank(method='first'), 4,  
      ↪labels=[1,2,3,4]).astype(int)  
rfm['m_quartile'] = pd.qcut(rfm['monetary'], 4, labels=[1,2,3,4]).astype(int)  
rfm['rfm_score'] = rfm['r_quartile'].astype(str) + rfm['f_quartile'].  
      ↪astype(str) + rfm['m_quartile'].astype(str)  
rfm.head()
```


```
[20]:
```

	customer_id	recency_days	frequency	monetary
r_quartile	f_quartile	m_quartile	rfm_score	
0	00012a2ce6f8dcda20d059ce98491703	288	1	114.74
2	1	3	213	

1	000161a058600d5901f007fab4c27140	410	1	67.41
1	1	2	112	
2	0001fd6190edaaf884bcaf3d49edf079	548	1	195.42
1	1	4	114	
3	0002414f95344307404f0ace7a26f1d5	379	1	179.35
1	1	4	114	
4	000379cdec625522490c315e70c7a9fb	150	1	107.01
3	1	3	313	

```
[22]: df_sales.to_csv("olist_cleaned_orders.csv", index=False)
      rfm.to_csv("olist_customer_rfm.csv", index=False)
```

```
plt.figure(figsize=(6,4))
plt.bar(top_products['product_id'].astype(str).head(10),
        top_products['revenue'].head(10))
plt.title("Top 10 products by revenue (example)")
plt.xticks(rotation=45)
plt.tight_layout()
plt.savefig("plots_top_products.png")
plt.show()
```



```
[ ]:
```